

# Laboratory - Deep Reinforcement Learning: Foundations and Practical Environment Setup for Real-World Applications

Franco Terranova

Université de Lorraine, CNRS, Inria, LORIA

*franco.terranova@inria.fr*

August 20, 2024



<https://terranovalfr.github.io/teaching/2024-EASSS-Course>

- 1 DRL Agents - Practical Setup
  - StableBaselines3
- 2 Environment - Practical Setup
  - OpenAI Gym
- 3 Grid World
  - POMDP
  - Matrix Representation
  - Generalization

## Components:

- **Libraries for Well-Established Algorithms:**

- **Stable Baselines3:**

- <https://stable-baselines3.readthedocs.io/en/master/>

- **Ray Rllib:** <https://docs.ray.io/en/latest/rllib/index.html>

- **TF-Agents:** <https://www.tensorflow.org/agents>

- **Keras-RL:** <https://github.com/keras-rl/keras-rl>

- **Neural Network Architecture:**

- Number of layers and neurons per layer
  - Activation functions (e.g., ReLU, Tanh)
  - Network type (e.g., feedforward, convolutional, recurrent)

- **Algorithm Hyperparameters:**

- Learning rate
- Batch size
- Discount factor ( $\gamma$ )
- Exploration strategy (e.g.,  $\epsilon$ -greedy for the DQN)
- ...


- **Optimizer Selection:**

- Adam
- Stochastic Gradient Descent
- ...

- **Additional Considerations (RL Specific):**

- **Reward Shaping:** Modify rewards to guide the agent towards desired behaviors
- **Training Stability:** Use techniques to enhance training stability
- ...

# Stable Baselines3



🏠 / DQN

[Edit on GitHub](#)

## DQN

Deep Q Network (DQN) builds on [Fitted Q-Iteration \(FQI\)](#) and make use of different tricks to stabilize the learning with neural networks: it uses a replay buffer, a target network and gradient clipping.

### Available Policies

|                               |   |
|-------------------------------|---|
| <code>MlpPolicy</code>        | alias of <code>DQNPolicy</code>                             |
| <code>CnnPolicy</code>        | Policy class for DQN when using images as input.            |
| <code>MultiInputPolicy</code> | Policy class for DQN when using dict observations as input. |

## Notes

- Original paper: <https://arxiv.org/abs/1312.5602>
- Further reference: <https://www.nature.com/articles/nature14236>

### Note

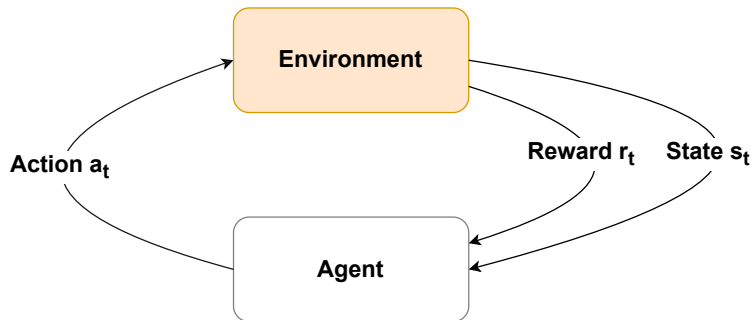
This implementation provides only vanilla Deep Q-Learning and has no extensions such as Double-DQN, Dueling-DQN and Prioritized Experience Replay.

## Can I use?

- Recurrent policies: ❌
- Multi processing: ✔️
- Gym spaces:

Website: <https://stable-baselines3.readthedocs.io/en/master/index.html>

# Environment Focus



- The environment maps an action and state to the next state and reward:

$$\text{Next State, Reward} = \mathcal{E}(\text{State, Action})$$

where  $\mathcal{E}$  represents the environment's dynamics

- Define internal dynamics so that agent can learn from it

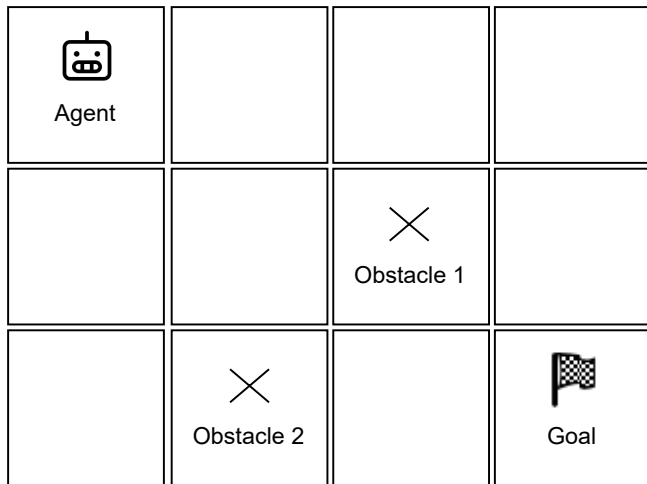


- **Standard Library:** Widely used for creating and testing RL environments
- **Class-Based:** Provides a Python class with attributes and methods to define and manage environments
- **Attributes:** Includes state and action spaces, attributes for rewards calculation, ...
- **Methods:** Simulate the dynamics and should adhere to the standard
- **Gymnasium:** Currently maintained version of Gym

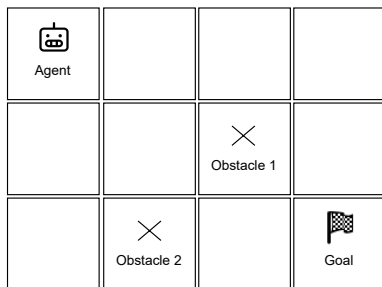
# OpenAI Gymnasium Methods

- `__init__`: Initializes the environment, setting up its initial dynamics and attributes
- `reset`: Resets the environment and returns:
  - **State**: The initial state
  - **Info**: A dictionary with additional information (optional)
- `step(action)`: Takes an action, updates the environment, and returns a tuple containing:
  - **Next State**: The state after the action
  - **Reward**: The reward received after taking the action in the state
  - **Done**: A boolean indicating if the episode has ended
  - **Truncated**: A boolean indicating if the episode has been truncated
  - **Info**: A dictionary with additional information (optional)
- `render`: Displays a visual representation of the environment (optional)
- `close`: Cleans up and closes the environment when done (optional)

# Grid World



# Grid World - Cell View



- **Observation:**  $(x, y)$  — The current position in the grid
- **Action:** Movement directions — up, down, left, right
- **Reward:**
  - Small Penalty — If the agent moves to an empty cell
  - Bigger Penalty — If the agent moves into an obstacle or outside
  - Prize — If the agent reaches the goal
- **Episode:** Terminates when the goal or cut-off is reached

- **Static Environment:**

- The environment does not change over time
- Finding the optimal policy is straightforward
- The cell view is a MDP for the task

- **General Case:**

- The grid evolves or changes
- Application to another grid with different parameters
- A generalizable policy may be complex to determine with this observation, being a POMDP for the task

- **Challenges on using this POMDP:**

- Same state leading to different action outcome
- Exposing multiple training environments may lead to instable learning

- **Full visibility:**

- Grid is translated into a 1D array
- Agent, Obstacles, Goal: Encoding choice
  - E.g. Current position represented with 1, Obstacles represented with 2s, and the goal with 3

- Now each observation will have a deterministic (reward, next state) when selecting an action

- **Assumption:** Fixed width and height for now

- **Periodic Changes:**

- The grid environment can be periodically switched (updated) with another version
- New environments can have different obstacle positions, and goal locations

- **Dynamic Training:**

- During training switch periodically based on a switch interval
- Ensures the agent can generalize across different environments


- What is the impact of ... ?
  - Episode cut-off
  - Rewards and their scale
  - Value of  $\gamma$
  - Number of training iterations
  - Relation between grid size and number of iterations
  - Epsilon decay
  - Number of environments



# Generalize across grid sizes

- **Grid size** determines the number of input and output neurons of the agent's NN
- Agent's NN specialized to a given grid size
- Possible solutions:
  - Re-train a NN for every grid size
  - Padding techniques to the maximum grid size
    - Need to set a maximum
    - Waste of resources for small grids
    - May require a larger NN

# Other solutions



- 1 Trade-off observability range
  - E.g. surrounding pixels in the 1-hop neighborhood
- 2 Features describing the environment
  - Requires manual features engineering
- 3 Potential combination of these solutions

# Convolutional Neural Networks

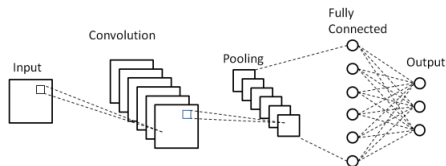


Image Source: Mesuga, Reymond Bayanay, Brian. (2021). A Deep Transfer Learning Approach to Identifying Glitch Wave-form in Gravitational Wave Data.

- Avoid feature engineering
- CNN will automatically determine a feature representation
- Updates driven by the DRL loss

# Well-Established Environments

## Popular Gym Environments:

- `CartPole-v1`: Balancing a pole on a moving cart
- `MountainCar-v0`: Driving a car up a hill
- `LunarLander-v2`: Landing a spacecraft on the moon
- Atari environments: Classic arcade games (e.g., `Pong-v0`, `Breakout-v0`)

## Example of Utilization:

```
import gymnasium as gym  
env_cartpole = gym.make('CartPole-v1')
```

- **State and Action Space Encodings:** Efficient representation that allows learning
- **Markovian Property:** Ensures that observation encodes all information needed
- **Exploration vs. Exploitation:** Balancing the trade-off between exploring new strategies and exploiting known ones
- **Sample Efficiency:** Improving the efficiency of learning algorithms to require fewer samples

- Addressing catastrophic forgetting with **continual reinforcement learning**
- Exploring scenarios with multiple agents (**Multi-agent RL**) and incorporating game theory
  - Independent learners
  - Cooperation games
  - Competitive games
- **Inverse Reinforcement Learning** to derive realistic reward functions
- **Meta Reinforcement Learning** enables fast learning across different tasks

# References



Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann (2021)

Stable-Baselines3: Reliable Reinforcement Learning Implementations.

*Journal of Machine Learning Research*, 22, 268:1-8.

<http://jmlr.org/papers/v22/20-1364.html>.



Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba (2016)

OpenAI Gym.

*arXiv:1606.01540*, <https://arxiv.org/abs/1606.01540>.



Yann LeCun, Yoshua Bengio, and Geoffrey Hinton (2015)

Deep learning.

*Nature*, 521(7553), 436–444.

# Advance Your RL Agents to New Horizons!